

Bayesian-Optimal Multi-Classification with Noisy Input Necessitates General Input Space Representation

Aman Bhargava

December 20, 2023

1 Introduction

Notation: lower case variables denote scalars (e.g., x), upper case variables denote random variables (e.g., X), and boldfaced variables denote vector quantities (e.g., \mathbf{x}, \mathbf{X}).

Here I analyze the latent representations in optimal Bayesian filter models trained to perform multi-class classification on some ground truth input vector $\mathbf{x}^* = [x_1^*, x_2^*]^\top$ based on noisy discrete-time measurement signals $\mathbf{X}(t) = [X_1(t), X_2(t)]^\top$ defined as

$$X_1(t) = x_1^* + \eta \mathcal{N}(0, 1) \quad (1)$$

$$X_2(t) = x_2^* + \eta \mathcal{N}(0, 1) \quad (2)$$

We scope our analysis to N linear classification boundaries. The multi-classification task may be defined in terms of N classification boundary angles $\alpha_1, \dots, \alpha_N$ which we collectively denote as Θ :

$$\Theta = \{\alpha_i : \alpha_i \in [-\pi, \pi], i = 1, \dots, N\} \quad (3)$$

as in Figure 1. Models are tasked with predicting the binary classification label for corresponding to each boundary in Θ . We denote the N labels $\mathbf{y}^* = [y_1^* \dots y_N^*]^\top$ corresponding to some \mathbf{x}^* based on the resulting noisy




Figure 1: **Multitasking RNN learns abstract representations.** Data generating process. The task is to simultaneously report whether the true joint evidence (x_1, x_2) (yellow dot) lies above (+1) or below (−1) a number of classification lines (here 6).

measurement signals $\mathbf{X}(t)$ where

$$y_i^* = \begin{cases} +1 & \text{if } x_2^* > x_1^* \tan \alpha_i \\ -1 & \text{otherwise} \end{cases} \quad (4)$$

Contribution: In this document, I demonstrate that a Bayesian-optimal multi-classifier with noisy random input $\mathbf{X}(t)$ and classification estimate output $\hat{\mathbf{y}}(t) \in [-1, 1]^N$ must form latent representations $\mathbf{z}(t)$ that retain the 2-dimensional structural information of the input data space $[x_1^*, x_2^*] \in \mathbb{R}^2$ in the limit as $N \rightarrow \infty$ for densely packed $\alpha_i \in \Theta$. Intriguingly, without noise, we are unable to guarantee that optimal latent representations $\mathbf{z}(t)$ will retain sufficient information to estimate \mathbf{x}^* . While the proof is stated for 2-dimensional input, the argument holds for any input dimensionality.

1.1 Bayesian Filtering Framework

Bayesian filters are a class of statistical models and algorithm that update a latent state based on noisy and uncertain observation signals. Rooted in principles of Bayesian inference, these filters combine aggregated “knowledge”, represented by a latent state $\mathbf{Z}(t)$, with incoming observations $\mathbf{X}(t)$ to continually update the latent state to facilitate some prediction of some output $\mathbf{Y}(t) = f(\mathbf{Z}(t))$.

Definition 1 (Bayesian Filter Operation). *A discrete-time Bayesian filter updates latent variable $\mathbf{z}(t)$ based on incoming data $\mathbf{x}(t)$ by applying Bayes’ theorem:*

$$P(\mathbf{z}(t)|\mathbf{x}(t), \mathbf{z}(t-1)) = \frac{P(\mathbf{x}(t)|\mathbf{z}(t), \mathbf{z}(t-1))P(\mathbf{z}(t)|\mathbf{z}(t-1))}{P(\mathbf{x}(t)|\mathbf{z}(t-1))} \quad (5)$$

$$\propto P(\mathbf{x}(t)|\mathbf{z}(t))P(\mathbf{z}(t)|\mathbf{z}(t-1)) \quad (6)$$

Bayesian filters are commonly equipped with a “decoder” or “readout map” f which maps latent $\mathbf{Z}(t)$ to readout estimation $\hat{\mathbf{Y}}(t) = f(\mathbf{Z}(t))$.

There is a deep structural similarity between RNNs and Bayesian filters, as both models update some latent state $\mathbf{z}(t)$ based on incoming datum $\mathbf{x}(t)$. Moreover, RNNs and Bayesian filters are both frequently used to predict some value $\mathbf{y}(t) = f(\mathbf{z}(t))$ (citation: Goodfellow for RNN, Bayesian inference textbook for filters). We leverage the structure in the Bayesian filter formulation to prove our main result in Section 2.

1.2 Canonical Bayesian Filtering Example

A canonical example of Bayesian filtering is estimating the position of a robot based on noisy sensory readings (Figure 2). Continuing with the same notation as above, assume the ground truth position of the robot is $\mathbf{x}^* \in \mathbb{R}^2$. We obtain noisy sensor readings $\mathbf{X}(t) = \mathbf{x}^* + \eta\mathcal{N}(\mathbf{0}, \mathbb{I})$, where η is the amount of zero-mean, identity covariance noise. The latent variable $\mathbf{Z}(t)$ represents our estimate of \mathbf{x}^* at time t . Using Equation 5, we obtain the updated estimated position based on incoming readings $\mathbf{X}(t)$.

In practice, we would initialize the latent variable $\mathbf{Z}(0) = \mathbf{0}$ or some other neutral position. The prior distribution $P(\mathbf{Z}(t)|\mathbf{Z}(t-1))$ represents our belief about the robot’s current position before seeing any new sensor data. We assume this is Gaussian-distributed around $\mathbf{Z}(t-1)$, though we can augment the sophistication of our prior model by including velocity and acceleration estimates (cf. Kalman filter). This added sophistication is how one may connect system state estimation (observer theory) to classic linear state space control theory. The likelihood function $P(\mathbf{X}(t)|\mathbf{Z}(t))$ represents the probability of observing a particular sensor reading given the current latent state $\mathbf{Z}(t)$. We assume this is Gaussian-distributed around $\mathbf{Z}(t)$. At each update

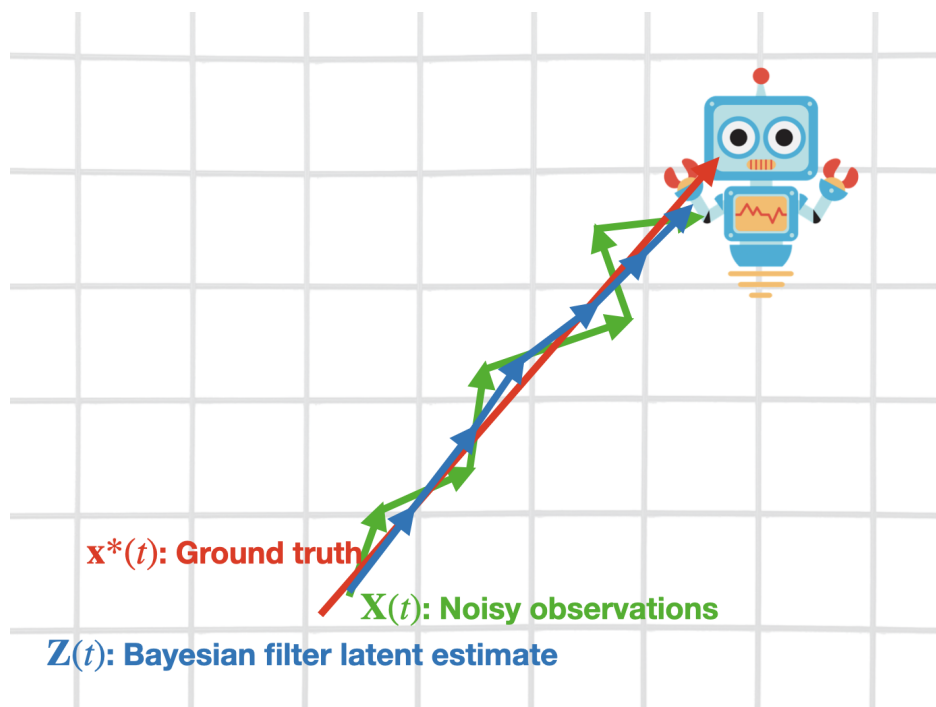


Figure 2: This figure shows the estimation of ground truth trajectory $\mathbf{x}^*(t)$ based on noisy readings $\mathbf{X}(t)$. $\mathbf{Z}(t)$ is shown in blue, representing the filtered estimate of $\mathbf{x}^*(t)$ in red based on $\mathbf{X}(t)$ in green.

step, the position estimate $\mathbf{Z}(t)$ is updated with the incoming data by combining the prior belief and the likelihood.

In neuroscience, this example parallels how the brain integrates information over time to refine its predictions over future states. This model highlights the probabilistic nature of perception and decision-making in the brain, illustrating how biological systems might deal with uncertainty and noise in sensory processing. Analyzing Bayesian-optimal filters is therefore of great interest for understanding the properties of high-performance iterative information processing systems – particularly those performing prediction tasks on uncertain/noisy input.

2 Main Results

Consider an optimal Bayesian filter for the multi-class classification task Θ on noisy discrete time measurement signals. Let ϵ be the maximum angular gap between classification boundaries in $\Theta = (\alpha_1 \dots \alpha_n)$ (Definition 7).

$$\epsilon = \max_{i \in N} \min_{j > i} \|\alpha_i - \alpha_j\| \quad (7)$$

Theorem 1. *An optimal Bayesian filter trained to perform multi-class classification on ground truth input \mathbf{x}^* w.r.t. decision boundaries $\Theta = \{\alpha_1, \dots, \alpha_N\}$ based on noisy measurement signals $\mathbf{X}(t)$ must have latent state $\mathbf{Z}(t)$ that retains a representation of the 2-dimensional input vector \mathbf{x}^* in the limit as $\epsilon \rightarrow 0$.*

Proof.

Lemma 1 (Equivalence to Angle Estimation). *In the limit as $\epsilon \rightarrow 0$ for decision boundaries $\Theta = (\alpha_1, \dots, \alpha_N)$, the multi-classification task of estimating \mathbf{y}^* (Equation 3) for a given \mathbf{x}^* given noisy observations $\mathbf{X}(t)$ is equivalent to estimating the angle $\theta = \angle \mathbf{x}^*$.*

Lemma 2 (Angle Estimation Requires Magnitude Estimation). *An optimal Bayesian filter predicting the angle of some ground truth input $\angle \mathbf{x}^*$ based on noisy observations $\mathbf{X}(t)$ must implicitly estimate the magnitude of \mathbf{x}^* during state updates on latent variable $\mathbf{Z}(t)$.*

Proof. We denote the conditional entropy of angle estimate $\hat{\theta} = f(\mathbf{Z}(t))$ as $H(\hat{\theta}|\mathbf{Z}(t)) = H(\hat{\theta}|\mathbf{X}(1), \dots, \mathbf{X}(t))$. Since $\mathbf{X}(t)$ is subject to equivariant Gaussian noise with variance η , the condition entropy is inversely proportional to the distance between point x^* and classification boundaries $\Theta = (\alpha_1, \dots, \alpha_N)$. For fixed angle $\theta = \angle \mathbf{x}^*$, the distance between \mathbf{x}^* and each classification boundary scales monotonically with $\|\mathbf{x}^*\|$. Therefore, the angle $\angle \mathbf{x}^*$ and the entropy of the angle estimate $H(\hat{\theta}|\mathbf{X}(t) \dots \mathbf{X}(t)) = H(\hat{\theta}|\mathbf{Z}(t))$ is sufficient to determine $\|\mathbf{x}^*\|$. Since $\hat{\theta}$ and $H(\hat{\theta}|\mathbf{Z}(t))$ are both functions of $\mathbf{Z}(t)$ and $\lim_{t \rightarrow \infty} \hat{\theta} = \theta$ we have demonstrated that $\mathbf{Z}(t)$ implicitly estimates the magnitude of \mathbf{x}^* . \square

Therefore, multi-class classification in the limit as the decision boundary spacing $\epsilon \rightarrow 0$ is equivalent to estimating the angle of the ground truth \mathbf{x}^* based on noisy $\mathbf{X}(t)$ (Lemma 1). Optimal Bayesian filtering to estimate the angle $\theta = \angle \mathbf{x}^*$ also implies estimating the magnitude of the ground truth data $\|\mathbf{x}^*\|$.

Thus we have demonstrated that both the angle information and the magnitude information of \mathbf{x}^* must be implicitly estimated by an optimal Bayesian filter in multi-class classification problem Θ with non-zero equivariant noise η in the measurements $\mathbf{X}(t)$. \square

3 Discussion

3.1 Implications of Non-Equivariant Gaussian Noise

In our analysis, we assumed equivariant Gaussian noise with variance η in $\mathbf{X}(t)$. The case of $\eta = [\eta_1, \eta_2]^\top$ with $\eta_1 \neq \eta_2$ can be made equivalent to the equivariant case via a coordinate transformation $x_2 = x_2\eta_1/\eta_2$. This may change the maximum inter-classification boundary distance ϵ , but the final result remains the same as $\epsilon \rightarrow 0$.

3.2 Generalizing to Higher Dimensions

To generalize these results to higher dimensions, the Lemmas 1 and 2 must be expanded. Lemma 1 focuses on the equivalence of the multi-classification problem to angle estimation. This condition can be expressed for a high dimensional space in terms of hyperplane decision boundaries passing through

the origin. The angle estimation problem becomes a multi-angle estimation problem (e.g., latitude and longitude for \mathbb{R}^3). The condition on maximal angular spacing $\epsilon \rightarrow 0$ can be generalized as an upper bound on surface area of the slices of a sphere defined by the hyperplanes' intersection with a sphere centered on the origin.