

Bayesian-Optimal Multi-Classification with Noisy Input Necessitates General Input Space Representation

Aman Bhargava

December 1, 2023

1 Introduction

Here we analyze the latent representations in Bayesian filter models trained to perform multi-class classification on some ground truth input vector $\mathbf{x}^* = [x_1^*, x_2^*]^\top$ based on noisy discrete-time measurement signals $\mathbf{X}(t) = [X_1(t), X_2(t)]^\top$ defined as

$$X_1(t) = x_1^* + \eta\mathcal{N}(0, 1) \quad (1)$$

$$X_2(t) = x_2^* + \eta\mathcal{N}(0, 1) \quad (2)$$

The multi-classification task may be defined in terms of N classification boundary angles which we collectively denote as Θ :

$$\Theta = \{\alpha_i : \alpha_i \in [-\pi, \pi], i = 1, \dots, N\} \quad (3)$$

as in Figure 1. Thus models must predict the ground truth label $\mathbf{y}^* = [y_1^* \dots y_N^*]^\top$ corresponding to some \mathbf{x}^* based on the resulting noisy measurement signals $\mathbf{X}(t)$ where

$$y_i^* = \begin{cases} +1 & \text{if } x_2^* > x_1^* \tan \alpha_i \\ -1 & \text{otherwise} \end{cases} \quad (4)$$




Figure 1: **Multitasking RNN learns abstract representations.** Data generating process. The task is to simultaneously report whether the true joint evidence (x_1, x_2) (yellow dot) lies above (+1) or below (−1) a number of classification lines (here 6).

Contribution: In this document, I demonstrate that a Bayesian-optimal multi-classifier with noisy random input $\mathbf{X}(t)$ and classification estimate output $\hat{\mathbf{y}}(t) \in [-1, 1]^N$ must form latent representations $\mathbf{z}(t)$ that retain the 2-dimensional structural information of the input data space $[x_1^*, x_2^*] \in \mathbb{R}^2$ in the limit as $N \rightarrow \infty$ for evenly spaced $\alpha_i \in \Theta$. Intriguingly, without noise, we are unable to guarantee that optimal latent representations $\mathbf{z}(t)$ will retain sufficient information to estimate \mathbf{x}^* . While the proof is stated for 2-dimensional input, the argument should hold for any input dimensionality.

1.1 Bayesian Filtering Framework

Bayesian filters are a class of statistical models and algorithm that update a latent state based on noisy and uncertain observation signals. Rooted in principles of Bayesian inference, these filters combine aggregated “knowledge”, represented by a latent state $\mathbf{Z}(t)$, with incoming observations $\mathbf{X}(t)$ to continually update the latent state to facilitate some prediction of some output $\mathbf{Y}(t) = f(\mathbf{Z}(t))$.

Definition 1 (Bayesian Filter Operation). *A discrete-time Bayesian filter updates latent variable $\mathbf{z}(t)$ based on incoming data $\mathbf{x}(t)$ by applying Bayes’*

theorem:

$$P(\mathbf{z}(t)|\mathbf{x}(t), \mathbf{z}(t-1)) = \frac{P(\mathbf{x}(t)|\mathbf{z}(t), \mathbf{z}(t-1))P(\mathbf{z}(t)|\mathbf{z}(t-1))}{P(\mathbf{x}(t)|\mathbf{z}(t-1))} \quad (5)$$

$$\propto P(\mathbf{x}(t)|\mathbf{z}(t))P(\mathbf{z}(t)|\mathbf{z}(t-1)) \quad (6)$$

Bayesian filters are commonly equipped with a “decoder” or “readout map” f which maps latent $\mathbf{Z}(t)$ to readout estimation $\hat{\mathbf{Y}}(t) = f(\mathbf{Z}(t))$.

There is a deep structural similarity between RNNs and Bayesian filters, as both models update some latent state $\mathbf{z}(t)$ based on incoming datum $\mathbf{x}(t)$. Moreover, RNNs and Bayesian filters are both frequently used to predict some value $\mathbf{y}(t) = f(\mathbf{z}(t))$ (citation: Goodfellow for RNN, Bayesian inference textbook for filters). We leverage the structure in the Bayesian filter formulation to prove our main result in Section 2.

2 Main Results

Consider an optimal Bayesian filter for the multi-class classification task on noisy discrete time measurement signals.

Theorem 1. *An optimal Bayesian filter trained to perform multi-class classification on ground truth input \mathbf{x}^* w.r.t. decision boundaries $\Theta = \{\alpha_1, \dots, \alpha_N\}$ based on noisy measurement signals $\mathbf{X}(t)$ must have latent state $\mathbf{Z}(t)$ that retains a representation of the 2-dimensional input vector \mathbf{x}^* in the limit as $N \rightarrow \infty$.*

Proof.

Lemma 1 (Equivalence to Angle Estimation). *In the limit as $N \rightarrow \infty$ for uniformly distributed decision boundaries in $\Theta = \{\alpha_i\}_{i \in [N]}$, the multi-classification task of estimating \mathbf{y}^* (Equation 3) for a given \mathbf{x}^* given noisy observations $\mathbf{X}(t)$ is equivalent to estimating the angle $\theta = \angle \mathbf{x}^*$.*

Lemma 2.

□

3 Conclusion